

# Certify-or-Refuse: A Machine-Checked Soundness Boundary for Untrusted Agent Edits to Opaque-Semantics Artifacts

2026-07-09

## Abstract

An LLM agent that edits a spreadsheet, a database schema, or a notebook produces a file that *opens fine* but may be silently wrong: a structural edit that fails to propagate references corrupts computed values with no visible symptom, and neither the agent nor a human reviewer can see it without the defining engine. We make the correctness of such an edit **checkable** rather than trusted. Our contribution is a formally verified soundness theorem, machine-checked in Lean 4 with no sorry and only the axioms `propext` and `Quot.sound`: evaluation of a computation is invariant under a function-and-dependency-preserving isomorphism, so a structural (coordinate-relabeling) edit whose reference-dependency graph is isomorphic to the original's reproduces every computed value — *under any semantics, without running the engine*. A companion locality theorem bounds what a value edit can affect to its downstream cone, and a third machine-checked theorem shows the tool's reference-shift *constructively produces* that isomorphism on a token-level formula model (single-cell and range-endpoint references) — narrowing the trusted base rather than eliminating it. On this spine we build a **certify-or-refuse router**: an untrusted agent's structural edit is accepted only when it equals the tool's own proven coordinate-shift transform, and otherwise explicitly refused — never silently wrong. We report a production certifier (`xlq certify`), the trusted reference-shift tokenizer it rests on (hardened to an exact grid-validity predicate and corroborated by value-preservation against an independent engine over 264 formulas), and — for the non-cell references a cell diff cannot see — a compare-and-fail-close net (defined-name and merge/hyperlink/autoFilter targets compared against the transform; data-validation, conditional-formatting, charts, pivots, external links refused) that we are careful to label an *enumerated denylist*, not a proven-complete allowlist, having ourselves shipped and then closed one silent false certification exactly where the enumeration was incomplete. We are deliberate about what is *proof* and what is *corroboration*: the theorem is the result; the empirical harness — an engine-free foreign-edit certifier that refuses 147/147 corrupted edits, an independent-oracle A/B in which the naive edit path silently corrupts 86.6% of real workbooks while the certified path corrupts none, and a diverse-corruptor confusion matrix in which every injected corruption is refused — corroborates it, and we report each with its confound (including a false certification our own review found

and we then closed), not as an independent soundness argument.

## 1. Introduction

The scarce resource in the agentic era is not capability but **verifiability**. A capable model will, with high probability, perform a structural spreadsheet edit correctly; but “high probability” is exactly the wrong guarantee for a financial model, a payroll sheet, or a regulatory filing, where a single unpropagated reference silently changes a computed total and ships. The failure is invisible by construction: the edited file is a valid `.xlsx`, it opens, and the wrong numbers sit in cells whose formulas *look* plausible. You cannot see the corruption without the engine, and asking the engine is asking the very component whose edit you are trying to check.

This paper asks a narrower, more durable question than “can the model do the edit?”: **can we certify that a given edit is value-faithful, offline, against the artifact’s own structure, without trusting the editor and without running the defining engine?** The answer we prove is yes for the *structural* fragment — the coordinate-relabeling edits (insert/delete rows and columns) that are simultaneously the most common agent operations and the ones whose damage is most invisible. The certifier accepts a structural edit only when its reference-dependency graph is the relabeling of the original’s, and refuses otherwise. Because the guarantee is grounded in the artifact rather than the model, it does not decay as models improve — it is the boundary a capable agent should be *required* to pass, not a crutch for a weak one.

Our claims, in order of strength:

1. **(Proof.)** A machine-checked theorem (Lean 4, no sorry) that value-faithfulness of a structural edit reduces to graph isomorphism of its reference-dependency structure — engine-free and semantics-agnostic — plus a locality theorem bounding value-edit effects to a downstream cone (§3).
2. **(Mechanism.)** A certify-or-refuse router and a production certifier that decide accept/refuse for an untrusted foreign edit by comparing it to the tool’s own proven transform, with a hardened trusted tokenizer and a fail-closed boundary for the one syntactically-undecidable case (§4).
3. **(Corroboration, reported with confounds.)** An engine-free foreign-edit certifier, an independent-oracle A/B, and an independent-oracle confusion matrix — each of which *supports* the theorem on real workbooks and each of which we report with its limitation, not as independent proof (§5).

A recurring, honest theme: three rounds of adversarial review of our own artifacts found real defects — silent-corruption bugs in the trusted tokenizer, a circular experiment, and an unimplemented soundness boundary — each of which we fixed and report as a fixed defect, because the discipline of finding them is part of the contribution (§6).

## 2. The problem: opaque-semantics artifacts and invisible structural damage

A spreadsheet is a program whose semantics live in an engine (Excel, LibreOffice Calc, IronCalc) that most tools do not reimplement. An edit tool that manipulates the file’s bytes therefore edits a program it cannot evaluate. The dominant failure mode is the **unpropagated reference**: insert a row above a formula’s input and, unless every reference below the insertion point is shifted, the formula now reads the wrong cell. The popular `openpyxl` library’s `insert_rows` does exactly this — it moves cell contents but rewrites zero references — so on any workbook with a below-insertion reference, it produces a file that opens cleanly and computes wrong (§5.2).

The artifact carries its own partial ground truth: Excel writes a cached value `<v>` next to each formula. This *self-oracle* is what a certifier can check against — but only for cells the artifact chose to cache, and only if the checker can recompute, which returns us to the engine. Our theorem removes the recompute: for the structural fragment, value-faithfulness follows from structure alone.

## 3. The formal core (the contribution)

We model a computation as  $C = (fn, deps)$ :  $fn\ n$  maps a node’s ordered dependency *values* to its value, and  $deps\ n$  is its ordered dependency *nodes*. This is exactly what a formula carries — what it reads and how it combines what it reads — with  $fn$  left an arbitrary (opaque) function. Evaluation is fuel-bounded:  $eval\ C\ 0\ \_ = default$  and  $eval\ C\ (k+1)\ n = fn\ n\ (map\ (eval\ C\ k)\ (deps\ n))$ . The theorem holds at every fuel level, hence for the true evaluation of any acyclic computation.

**Theorem 1 (exact structural certification).** Let  $\sigma$  map every node of  $C$  to a node of  $C'$  with the same function and dependency list relabeled by  $\sigma$  (a function- and edge-preserving graph isomorphism):  $C'.fn\ (\sigma\ n) = C.fn\ n$  and  $C'.deps\ (\sigma\ n) = map\ \sigma\ (C.deps\ n)$ . Then  $eval\ C'\ k\ (\sigma\ n) = eval\ C\ k\ n$  for all  $k, n$ .

*Consequence.* If the original’s evaluation is the artifact’s embedded ground truth  $0$  ( $eval\ C\ k = 0$ , the self-oracle), then a structurally-faithful edit reproduces that ground truth at the relabeled positions,  $eval\ C'\ k\ (\sigma\ n) = 0\ n$ , established without the defining engine and without recomputing  $0$ . This is the exact tier of certify-or-refuse.

**Theorem 2 (locality / audit-surface bound).** Define  $agree\_upto\ C\ C'\ k\ n$ : the two computations have identical  $fn$  and  $deps$  at every node within  $k$  dependency-steps of  $n$ . Then  $eval\ C\ k\ n = eval\ C'\ k\ n$  — a node’s value depends only on its dependency cone. Hence a value edit changes nothing outside its downstream cone: a mixed edit factors into a certified structural scaffold plus value fills whose effect is provably contained, collapsing the audit surface from “did this corrupt anything anywhere?” to “are these  $N$  cells right?”.

Both theorems are machine-checked in Lean 4, self-contained (no Mathlib), with `#print axioms` reporting only `[propext, Quot.sound]` and no sorry (`formal/SelfOracle.lean`). The reference-shift algebra’s arithmetic laws (`insert•delete = identity`, monotonicity,

the six-case delete clamp against the set-theoretic truth) are separately proved for all inputs by the Z3 SMT solver (formal/shift\_laws.py).

**Theorem 3 (graph preservation — the shift discharges the hypothesis, for single cells and range endpoints).** Theorems 1 and 2 *assume* the edited graph is the  $\sigma$ -relabeling of the original's; the tool does not assume it, it produces it by shifting references. We model a formula as a list of tokens — each an opaque literal (number, string, function name), a single cell reference, or a range carrying its two endpoints — and machine-check that the reference-shift produces exactly the relabeled graph:  $\text{refs}(\text{shiftF } \sigma \text{ } f) = (\text{refs } f).\text{map } \sigma$ , with literals provably untouched (formal/RefShift.lean). This is *verbatim* the hdeps premise of Theorem 1, so for these token shapes the graph isomorphism is discharged **constructively by the tool's operation**, not assumed. We also machine-check  $\text{delete} \circ \text{insert} = \text{id}$  on the cell maps. All sorry-free, axioms [propext, Quot.sound].

**What Theorem 3 does and does not give.** It is a fusion law parametric in an *arbitrary total* map  $\sigma : \text{Cell} \rightarrow \text{Cell}$ : it proves the shift *transports* the reference graph by whatever  $\sigma$  it is given — it does **not** verify that  $\sigma$  is Excel's shift arithmetic (that is the separate Z3-proved algebra), nor does the total-map model express the asymmetric six-case delete *clamp* ( $\text{head-in-band} \rightarrow k$ ,  $\text{tail-in-band} \rightarrow k-1$ ), the #REF! a fully-consumed reference produces, or the true multi-cell interior of a range (we model a range by its endpoints). So the honest statement is: **the token→graph→value chain is machine-checked for single-cell references and range-endpoint pairs under a value-preserving relabeling**; the explicitly *trusted* surface is the byte→token parse (§4, §5.1), the correctness of  $\sigma$  itself, the asymmetric clamp / #REF! algebra, range-interior dependencies, and sheet-qualified/3D/table/external references. This is a real narrowing of the trusted base — not its elimination. Fidelity of value edits remains out of the exact tier by construction.

## 4. The certify-or-refuse router and its trusted base

**The router.** Given an original artifact and an untrusted foreign edit plus a declared structural op, the router computes the op's node relabeling  $\sigma$ , and CERTIFIES iff the foreign edit's reference-dependency graph is exactly the  $\sigma$ -image of the original's; otherwise it REFUSES. Two properties make this sound as a check of *untrusted* work: under-declaring a change is caught (an unaccounted graph difference  $\rightarrow$  refuse), and a wrong  $\sigma$  yields a mismatch  $\rightarrow$  refuse. The router is fail-closed: any condition it cannot certify becomes a refusal.

**The production certifier.** `xlq certify <orig> <edited> --op ... --at ...` realizes the router with the tool's *complete* formula parser: it applies the tool's own proven structural transform to the original and diffs the result against the foreign edit, certifying iff the only differences are stripped/rewritten caches and number formats (which foreign tools routinely touch), and refusing on any formula, value, added, or removed-cell difference. It is engine-free (it compares stored formulas and raw data, never recomputed values) and multi-sheet-safe (it diffs the union of sheets). Certification therefore means “this foreign edit equals the tool's proven transform”; its marginal value over simply *doing* the edit is trust-topology — the untrusted producer's artifact is what

ships, gated fail-closed, with the tool off the write path — not any verifier-cheaper-than-prover asymmetry (there is none; the certifier recomputes the transform).

**The trusted base: the reference-shift tokenizer.** Every certificate is only as sound as the predicate deciding *which tokens in a formula are cell references*. We initially used syntactic proxies (a 1-3-letter column, an unbounded row) and adversarial review found these were wrong in three ways that caused *silent corruption*: a function name whose prefix looks like a cell (BIN2DEC  $\rightarrow$  BIN3DEC under a row insert), a function name ending in digits before a paren (LOG10  $\rightarrow$  LOG11), and out-of-grid tokens (XFE9, A2000000) whose column or row exceeds the sheet limits. We replaced the proxies with the exact **grid-validity** predicate — a token is a cell reference iff its column is in A..XFD (1..16384) and its row is in 1..1048576, boundary-gated so an identifier or function call is never mistaken for a reference. This one rule subsumes all four bug classes.

We validate the tokenizer at two levels. First, **differential cross-checking** against an independent A1 shifter over 19 digit-bearing Excel functions, out-of-grid tokens, \$-absolutes, and ranges, across insert-rows and delete-rows: 175 (formula, op) pairs, 0 disagreements (tokenizer\_fuzz.py). This establishes *mechanization consistency* between the Rust scanner and the reference — but the reference is a second implementation of *our own* grid-validity spec, so it cannot establish conformance to a spreadsheet engine’s semantics. Second, therefore, **differential value-preservation against an independent engine**: a blank-row insert is value-preserving under a correct reference shift, so recomputing the same file with LibreOffice *before and after* the edit must leave every formula’s value unchanged (tokenizer\_conformance.py, seeded property-based generation over evaluable formulas — digit-bearing function names, single/mixed/absolute refs, ranges, strictly-distinct cell values so a mis-shift onto another cell changes the value). We run this across **all four structural ops** (insert/delete  $\times$  rows/columns) and against **two independent engines** — LibreOffice *and* the pure-Python formulas engine, neither of which is xlq’s IronCalc (conformance\_v2.py): **465 formulas checked per engine, 0 value divergences, with the two engines agreeing exactly**. (The earlier single-op/one-engine run, tokenizer\_conformance.py, gave 264 formulas / 0 divergences.) We are precise about what this does and does not establish. Value-preservation is *necessary but not sufficient* for reference-graph correctness: a mis-shift of a reference that the formula’s value does not depend on (a zero-weighted term, a non-max argument of MAX) would pass undetected, so the count over-counts references actually *witnessed* by the value. Two independent engines agreeing rules out a single-engine artifact, but both are open-source engines, not Excel; and sheet-qualified/3D/table/R1C1 constructs are outside the generator. The honest label is *differential value-preservation against LibreOffice — a necessary-not-sufficient corroboration of the trusted parse, complementing the same-spec fuzzer* — not full conformance to an engine’s reference semantics.

**The one undecidable case, made a fail-closed boundary.** A defined name spelled like a grid-valid cell (FY2021 = column FY, row 2021; Q1; Tax2020) is indistinguishable from a reference *in the formula text*, so the tokenizer would shift its uses and the edited file would still equal the tool’s own (wrong) transform — the single place certified  $\Rightarrow$  correct could be false on a real workbook. It is, however, decidable from

the workbook’s defined-names table, which the certifier already holds. We close it fail-closed: a structural edit on a workbook containing a defined name that collides with a cell-reference spelling emits a `defined_name_ref_collision` residual, and the certifier refuses. A workbook defining FY2021 and using it in a formula routes to REFUSE; a benign name (TaxRate) proceeds. This converts the last hidden hole into a demonstrated boundary and scopes the verified surface explicitly: **single sheet, in-grid coordinates, row/column shifts, no name collision.**

## 5. Corroboration on real workbooks (reported with confounds)

None of the following is the soundness argument — that is §3. Each supports the theorem on real files and each carries a limitation we state.

### 5.1 Engine-free certification of untrusted foreign edits

Running the router (format-parametric core) on openpyxl’s edits of 172 real formula-bearing workbooks, engine-free, cross-checked against an independent LibreOffice oracle: **0 false certifications and 147/147 corrupted foreign edits refused** (`foreign_certify`). The load-bearing caveat, always reported alongside: this soundness is bought with heavy refusal — the fail-closed A1 proxy certifies only about **1 in 23 faithful** foreign edits, because it refuses every reference form it cannot fully model (whole-row, whole-column, cross-sheet, table, defined-name). Useful soundness — a high certify rate on faithful edits — requires the complete parser of §4, not the proxy. The value is the boundary, not the throughput.

### 5.2 Independent-oracle A/B: the corruption the boundary prevents

On 172 real workbooks, `insert-row@2`, with LibreOffice (independent of both openpyxl and the tool’s engine) recomputing each edit against Excel’s cache: the naive openpyxl edit path **silently corrupts 86.6%** of workbooks (149/172), while the tool’s structural edit is engine-confirmed faithful on 150/172 and explicitly refuses the remaining 22 — **0 silent corruptions** (`agent_ab`). We extend this across **all four structural ops and both independent engines** on generated distinct-value workbooks (`agent_ab_v2.py`): the naive path silently corrupts insert-rows 100% / delete-rows 98.8% / insert-cols 94.3% / delete-cols 80.2% (**93.8% overall**), the tool **0% on every op**, and the formulas engine agrees with LibreOffice *exactly* — so the protection is not a LibreOffice artifact and holds across the op space, not just `insert-row@2`. We flag the confound that remains: the openpyxl figure is a corpus/tool property (openpyxl shifts no references), not an agent-error distribution, and a competent reference-shifting engine also gets these ops right; the tool’s differentiator is auditability and explicit refusal, which the A/B does not isolate. We found and fixed an oracle false-positive during this study — `ROW()/COLUMN()` are position-dependent and legitimately change on a row insert — which is why the oracle excludes position-dependent functions.

On the **real** corpus across all four ops we check shift correctness *deterministically*

rather than by recompute, because real-corpus value-preservation is confounded — LibreOffice recomputes exotic financial/date functions (ACCRINT, CUMPRINC, DB, DAYS360, TIME) inconsistently with the Excel cache, flagging as “corrupt” files whose shifted formulas we verified correct by hand. So we compare xlq’s output formulas to the two-engine-validated reference shifter over  $\sim 6,000$  real formula cells (shift\_correctness\_real.py): **xlq is 100% correct on every op** (insert-rows 1651 cells, delete-rows 1608, insert-cols 1648, delete-cols 1076; 0 mismatches), while the naive path leaves 17-72% of shift-requiring cells wrong. Constructs the simple checker cannot independently verify (whole-column/row, cross-sheet, tables, range-with-function endpoint) are skipped, not guessed; every early “mismatch” was a checker bug (comparing a deleted-band cell to the wrong output cell; the reference shifter not handling F:F, which xlq correctly shifts to G:G), and xlq had zero real errors.

## 5.2b A fifth op — move-rows — shows $\sigma$ is not special to insert/delete

Theorem 1 holds for *any* function-and-dependency-preserving isomorphism  $\sigma$ , so the certifiable class is not limited to the monotonic insert/delete shift. We add **move-rows** (relocate a contiguous block of rows), whose  $\sigma$  is a *permutation* — proved a bijection in Z3 (injective + surjective on  $[1, \text{maxrow}]$ , both directions) — and whose physical edit reorders rows (a buffered, non-streaming rewrite). A range whose endpoints reorder under the permutation cannot be a shifted rectangle, so it fail-closes as a move\_straddles\_range residual. On the same real corpus, deterministic move-rows shift correctness is **100% on 1,538 real formula cells** (25/60 workbooks fail-closed as residual/straddle), and the certify-or-refuse contract holds: certify(orig, correct move) CERTIFIES while a botched move (one reference reverted) is REFUSED. This is a concrete demonstration that the exact tier extends to a genuinely non-monotonic structural edit by supplying the right  $\sigma$ , with no change to the theorem.

## 5.3 Independent-oracle confusion matrix for the production certifier

We adjudicate xlq certify’s verdicts *independently of the tool* over **diverse** corruption. An initial matrix used a single corruptor — openpyxl’s deterministic no-op-shift — and the Excel cache with a reliability gate; its FN=0 on 14 edits was a point estimate (rule-of-three  $\approx 21\%$ ) confounded with editor identity. We diversified to three corruptor types — openpyxl (no-op), unshift\_one (revert one shifted reference), and wrong\_delta (over-shift one reference) — with ground truth **by construction** (we inject the corruption, so its label is independent of any engine) and a LibreOffice *self-consistent* oracle as cross-check (recompute before/after; no reliability gate). All 45 injected corruptions were refused (each type 15/15) and all 15 faithful tool-produced edits certified (cert\_confusion\_v2). We are careful about what this does *not* show: because the certifier accepts iff the edit equals the tool’s transform, and all three corruptors are *defined* as deviations from that transform, FN=0 here is **by construction**, not a sampled bound — we do not attach a confidence interval to a structurally-guaranteed zero, and unshift\_one/wrong\_delta differ only in the sign of a perturbation the certifier does not weigh. The one *reachable* false certification is not in this

family at all: the cell diff compares only sheet cells, so a foreign edit that shifts every cell formula correctly but leaves a **non-cell reference** (a defined name’s target, a data-validation or conditional-formatting formula, a chart series) unshifted is invisible to it. We built exactly that edit (`cert_noncell_test.py`): with a defined name `Rate = $A$10` left unshifted while all cells moved, an earlier xlq certify **CERTIFIED it with all-zero diffs — a genuine false certification**. We closed it: certify now checks that defined names match the tool’s transform (refusing the unshifted one, certifying the correct one) and *fail-closes* on data-validation, conditional-formatting, and chart parts it does not compare — a stated coverage boundary. Separately, 6 of the cell-formula corruptions were **value-preserving** (a reference error that does not change the current value); the value oracle called them faithful but certify refused all 6 — its structural equality is *stricter* than a value check. The symmetric cost, stated plainly: certify equally refuses a *faithful but non-identical* rewrite (a commuted  $A1+B1 \rightarrow B1+A1$ ), so its accept class is “byte-for-formula identical to the tool’s transform,” narrower than “value-faithful.”

## 5.4 The interventional finding: differential testing hardened the trusted base

We ran a live fast-model agent on 20 real workbooks, tasking it to rewrite formula references for `insert-row@2`, and compared its output to the tool’s transform. As a guard-soundness experiment this is **circular** — the “agent correct” label and the xlq certify verdict are the same predicate applied to a file grafted onto the tool’s own output — and we do not report it as such. Its genuine, non-circular yield is **differential testing that surfaced real silent-corruption bugs in the tool’s own tokenizer** (BIN2DEC, Sales2020, and on follow-up LOG10 and the out-of-grid class), each now fixed and covered by §4’s cross-check. A validation method that finds real bugs in its own trusted base is exactly what a certifier’s TCB needs; the honest framing is that the method proved its worth, not that it produced a soundness number.

## 6. On finding our own defects

Successive adversarial reviews of these artifacts each landed a real hole: the tokenizer’s syntactic proxies (silent corruption); the circular live-agent experiment (a tautological “0 false certifications”); the unimplemented defined-name *aliasing* boundary; and — most consequentially — a **reachable false certification** in the production certifier itself, where a foreign edit that shifts every cell formula correctly but leaves a defined name’s target unshifted was certified with all-zero diffs, because the cell diff never compared non-cell references. We built the exploit, confirmed it, and closed it (defined names now compared against the tool’s transform; data-validation, conditional-formatting, charts, pivots, and external links fail-closed). A subsequent review then falsified a *universal* phrasing of that fix by the same pattern — a foreign edit that reverts a `mergeCell/hyperlink/autoFilter` reference, which the transform shifts but the cell diff never compared, was certified — so we extended the net to those semantic structural references and reframed the boundary as an explicitly *enumerated denylist* whose completeness we do not claim to have proven. We report each as a fixed defect. This is part of the contribution: a certify-or-refuse claim is only

credible if its authors have tried hardest to break it — the record of what broke, and what the fix was, is the evidence that the remaining boundary is real, and it is why we are conservative about calling the corroboration anything more than corroboration, and the denylist anything more than a denylist.

## 7. Scope and limitations

The **verified surface** the certifier certifies (rather than refuses) is: row/column insert/delete *and row-block move* (§5.2b), single-sheet, in-grid coordinates, cell formulas over single-cell and range-endpoint references, with no defined-name/cell collision and (for move) no move-straddling range. On top of the sheet-cell diff, certify explicitly compares the reference-bearing constructs the transform shifts that a cell diff cannot see — defined-name targets and mergeCell/hyperlink/autoFilter refs — refusing any that differ from the transform, and fail-closes on data-validation, conditional-formatting, sparklines, charts, pivot tables, and external links. Two honesty caveats, both learned from this artifact’s own failures. First, this compare-and-fail-close net is an **enumerated denylist**, not a machine-checked allowlist: its completeness over value-affecting non-cell references is *asserted, not proven*, so an unenumerated value-bearing construct would be silently certified — the exact failure the defined-name case exhibited until a reviewer built the exploit. (We keep certify’s compare surface a superset of the transform’s *value-bearing* write surface — every reference-bearing construct except the deliberately-excluded view-state dimension/selection/pane/brk — which makes the enumeration mechanically checkable against `has_ref_attr`, but that superset relation is a code invariant, not a proof.) Second, the guarantee is over **computed values**: pure view-state the transform also shifts — the used-range dimension, selection, frozen pane, and page brks — is deliberately *not* compared, because it is non-value-bearing and foreign tools legitimately vary it; a certified file may differ from the transform there without affecting any value. Value edits are out of the exact tier by construction. The **trusted base** (narrowed by Theorem 3 but not eliminated) is: the byte→token parse; the correctness of the shift map  $\sigma$  itself (the Z3-proved arithmetic, not re-proved in the graph layer); the asymmetric six-case delete clamp and its #REF! outcomes; the multi-cell interior of ranges (modeled by endpoints); the constructs absent from the token type; and the completeness of the non-cell denylist. The corroboration is bounded accordingly: value-preservation is necessary-not-sufficient (though now confirmed across all four structural ops and by two independent engines that agree, ruling out a single-engine artifact, and still not Excel); the confusion matrix’s FN=0 is by-construction over deviations-from-transform, so the informative test is the non-cell corruptor of §5.3, which found and closed a real false certification. The exact tier certifies a measured 37.5% of operations and 27% of whole tasks on a realistic edit-distribution study; the majority of real tasks are mixed and need the probabilistic tier — but by Theorem 2 a mixed task factors into a *certified structural scaffold* plus a bounded value-fill cone, and we measure this (`composition_coverage.py`): on mixed edits the scaffold is always certified and the audit surface collapses to the value-fill cone (100/84/68/52% of the artifact certified untouched as fills grow, mean 76%), so the certifiable *component* rises from **27% fully certified to 87% with a certified scaffold**. A verified byte-level tokenizer, and a model faithful to the clamp/#REF! algebra, are the natural

next steps.

## 8. Related work

Spreadsheet analysis and smell detection, structural-edit tools that reshape files without a fidelity property, and record/replay or unit-test approaches to spreadsheet correctness all differ from this work in the same way: they establish behavior on tested inputs, whereas we prove value-fidelity of the structural fragment for all inputs and all semantics, and gate untrusted edits on that proof. Verified compilation and translation validation are the closest methodological analogues — accept a producer’s output only when it matches a proven-correct reference — which we adapt to the setting of opaque-semantics artifacts edited by untrusted agents.

## 9. Conclusion

Verifiability, not capability, is the durable constraint on agent edits to opaque-semantics artifacts. We proved — machine-checked, engine-free, semantics-agnostic — that value-fidelity of a structural spreadsheet edit reduces to a graph-isomorphism check, built a certify-or-refuse router that gates untrusted edits on that proof with a fail-closed boundary for the one undecidable case, and corroborated it on real workbooks while stating each confound. The result ships as an honest, genuinely-verified boundary over an explicitly-scoped surface — the boundary a capable agent should be required to pass, not a patch for a weak one.